

## ✓ DS340 Final Project

### Image Classification of Luggage X-ray Dataset

For this project, we use this dataset from: <https://universe.roboflow.com/airport-security-scanning/airport-security-scans-dataset>.

As you can see in this data set we have images of X-rays of luggage and it is classified into 5 different types of contraband. Our goal is to create a model that can accurately classify these images into the correct type of contraband. If AI can be accurate enough to detect contraband in luggage it can make traveling much safer and give an extra set of eyes to TSA.

```
from google.colab import drive
import zipfile

zip_path = '/content/drive/MyDrive/ds340/project/data.zip'

with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall('/content/drive/MyDrive/ds340/project/')

print(os.listdir('/content/drive/MyDrive/ds340/project/'))
```

## ✓ Load the data

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

Mounted at /content/drive

!ls '/content/drive/MyDrive/ds340/project/'

data  data1  data_copy  data.zip  idk  __MACOSX  ogdata  originalData

!ls '/content/drive/MyDrive/ds340/project/data1/test'

'Folding Knife'  'Multi-tool Knife'  Scissor  'Straight Knife'  'Utility Knife'
```

```

import os

train_dir = '/content/drive/MyDrive/ds340/project/data1/train'
valid_dir = '/content/drive/MyDrive/ds340/project/data1/valid'
test_dir = '/content/drive/MyDrive/ds340/project/data1/test'
test = '/content/drive/MyDrive/ds340/project/data1/test'

print(os.listdir(test_dir))
print(os.listdir(test))

['Utility Knife', 'Folding Knife', 'Multi-tool Knife', 'Straight Knife', 'Scisso
['Utility Knife', 'Folding Knife', 'Multi-tool Knife', 'Straight Knife', 'Scisso

from google.colab import drive
import os

def count_files(directory):
    for subdir in os.listdir(directory):
        subdir_path = os.path.join(directory, subdir)
        if os.path.isdir(subdir_path):
            num_files = len(os.listdir(subdir_path))
            print(f'{subdir} contains {num_files} files')

print("Training Data:")
count_files(train_dir)

print("\nValidation Data:")
count_files(valid_dir)

print("\nTest Data:")
count_files(test_dir)

```

Training Data:

Utility Knife contains 1110 files  
 Folding Knife contains 1102 files  
 Multi-tool Knife contains 1122 files  
 Straight Knife contains 578 files  
 Scissor contains 1039 files

Validation Data:

Utility Knife contains 327 files  
 Folding Knife contains 319 files  
 Multi-tool Knife contains 330 files  
 Straight Knife contains 148 files  
 Scissor contains 291 files

Test Data:

Utility Knife contains 177 files  
 Folding Knife contains 155 files  
 Multi-tool Knife contains 149 files  
 Straight Knife contains 79 files  
 Scissor contains 147 files

## ▼ Visualize Data Distribution

```

import os
import matplotlib.pyplot as plt

base_dirs = {
    'train': train_dir,
    'valid': valid_dir,
    'test': test_dir
}

data_summary = {}
total_files = 0

for key, directory in base_dirs.items():
    data_summary[key] = {}
    for subdir in os.listdir(directory):
        subdir_path = os.path.join(directory, subdir)
        if os.path.isdir(subdir_path):
            num_files = len([name for name in os.listdir(subdir_path) if os.path.isfile(name)])
            data_summary[key][subdir] = num_files
            total_files += num_files

labels = list(data_summary['train'].keys())
train_counts = [data_summary['train'][label] for label in labels]
valid_counts = [data_summary['valid'][label] for label in labels]
test_counts = [data_summary['test'][label] for label in labels]

fig, axes = plt.subplots(1, 3, figsize=(18, 6), sharey=True)
fig.suptitle('Distribution of Image Types Across Folders')

axes[0].bar(labels, train_counts, color='b')
axes[0].set_title('Train Data')
axes[0].set_ylabel('Number of Images')
axes[0].set_xlabel('Image Type')

axes[1].bar(labels, valid_counts, color='g')
axes[1].set_title('Validation Data')
axes[1].set_xlabel('Image Type')


axes[2].bar(labels, test_counts, color='r')
axes[2].set_title('Test Data')
axes[2].set_xlabel('Image Type')

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

print("Percentage Distribution Across Folders:")
for label in labels:
    total = train_counts[labels.index(label)] + valid_counts[labels.index(label)] + test_counts[labels.index(label)]
    print(f"\n{label}:")
    print(f"  Train: {train_counts[labels.index(label)] / total * 100:.2f}%")
    print(f"  Valid: {valid_counts[labels.index(label)] / total * 100:.2f}%")
    print(f"  Test: {test_counts[labels.index(label)] / total * 100:.2f}%")

```

```
print(f" Test: {test_counts[labels.index(label)] / total * 100:.2f}%")
```



### Percentage Distribution Across Folders:

#### Utility Knife:

Train: 68.77%

Valid: 20.26%

Test: 10.97%

#### Folding Knife:

Train: 69.92%

Valid: 20.24%

Test: 9.84%

#### Multi-tool Knife:

Train: 70.08%

Valid: 20.61%

Test: 9.31%

#### Straight Knife:

Train: 71.80%

Valid: 18.39%

Test: 9.81%

#### Scissor:

Train: 70.35%

Valid: 19.70%

Test: 9.95%

Double-click (or enter) to edit

## ✓ Attempt 1 Multiclass Classification using CNN

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import cv2
import os
os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2"
import warnings
warnings.filterwarnings('ignore')

from sklearn.metrics import confusion_matrix, classification_report

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Activation, BatchNormalization, Conv2D, Dense, Dropout
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import CategoricalCrossentropy
from tensorflow.keras.regularizers import l2
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping

IMG_WIDTH = 224
IMG_HEIGHT = 224
BATCH_SIZE = 32

train_datagen = ImageDataGenerator(rescale=1.0/255,
                                   zoom_range=0.2,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   fill_mode='nearest')
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    target_size=(IMG_WIDTH, IMG_HEIGHT),
                                                    batch_size=BATCH_SIZE,
                                                    class_mode='categorical',
                                                    shuffle=True)

Found 4951 images belonging to 5 classes.
```

```
validation_datagen = ImageDataGenerator(rescale=1.0/255)
validation_generator = validation_datagen.flow_from_directory(valid_dir,
                                                             target_size=(IMG_WIDTH,
                                                             batch_size=BATCH_SIZE,
                                                             class_mode='categorical
                                                             shuffle=True)
```

Found 1415 images belonging to 5 classes.

```
labels = {value: key for key, value in train_generator.class_indices.items()}
```

```
print("Label Mappings for classes present in the training and validation datasets\n"
      for key, value in labels.items():
        print(f"{key} : {value}")
```

Label Mappings for classes present in the training and validation datasets

```
0 : Folding Knife
1 : Multi-tool Knife
2 : Scissor
3 : Straight Knife
4 : Utility Knife
```

```
x_batch, y_batch = train_generator.next()
print(x_batch.shape)
print(y_batch.shape)
```

```
(32, 224, 224, 3)
(32, 5)
```


```
fig, ax = plt.subplots(nrows=2, ncols=5, figsize=(15, 12))
idx = 0
```

```
for i in range(2):
    for j in range(5):
        label = labels[np.argmax(train_generator[0][1][idx])]
        ax[i, j].set_title(f"{label}")
        ax[i, j].imshow(train_generator[0][0][idx][:, :, :])
        ax[i, j].axis("off")
        idx += 1
```


```
plt.tight_layout()
plt.suptitle("Sample Training Images", fontsize=21)
plt.show()
```

## Sample Training Images

Scissor



Straight Knife




Multi-tool Knife



Folding Knife




Utility Knife



Multi-tool Knife




Folding Knife



Folding Knife



Scissor



Folding Knife



✓ Model #1



```

train_datagen = ImageDataGenerator(rescale=1./255)
valid_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical')

validation_generator = valid_datagen.flow_from_directory(
    valid_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical')

```

Found 4951 images belonging to 5 classes.  
 Found 1415 images belonging to 5 classes.

```

from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

```

```

model = Sequential([

    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(2, 2),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),

    Dense(512, activation='relu'),
    Dense(5, activation='softmax')
])

```

```
model.compile
```

#### **keras.src.engine.training.Model.compile**

```
def compile(optimizer='rmsprop', loss=None, metrics=None, loss_weights=None,
weighted_metrics=None, run_eagerly=None, steps_per_execution=None,
jit_compile=None, pss_evaluation_shards=0, **kwargs)
```

turns on exact evaluation and uses a heuristic for the number of shards based on the number of workers. 0, meaning no visitation guarantee is provided. NOTE: Custom implementations of `Model.test\_step` will be ignored when doing exact evaluation. Defaults to `0`.

**\*\*kwargs:** Arguments supported for backwards compatibility only.

```
model.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_15 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_16 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_16 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_17 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_17 (MaxPooling2D)	(None, 17, 17, 128)	0
flatten_5 (Flatten)	(None, 36992)	0
dense_10 (Dense)	(None, 512)	18940416
dense_11 (Dense)	(None, 5)	2565
Total params: 19036229 (72.62 MB)		
Trainable params: 19036229 (72.62 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(
    train_generator,
    steps_per_epoch=100,
    epochs=15,
    validation_data=validation_generator,
    validation_steps=50
)
```

Epoch 1/15

100/100 [=====] - ETA: 0s - loss: 1.6692 - accuracy: 0.

100/100 [=====] - 85s 762ms/step - loss: 1.6692 - accuracy: 0.

Epoch 2/15

100/100 [=====] - 57s 570ms/step - loss: 1.5952 - accuracy: 0.

Epoch 3/15

100/100 [=====] - 58s 574ms/step - loss: 1.5849 - accuracy: 0.

Epoch 4/15

100/100 [=====] - 61s 603ms/step - loss: 1.5838 - accuracy: 0.

Epoch 5/15

```

100/100 [=====] - 61s 607ms/step - loss: 1.5910 - accur
Epoch 6/15
100/100 [=====] - 58s 572ms/step - loss: 1.5857 - accur
Epoch 7/15
100/100 [=====] - 65s 645ms/step - loss: 1.5838 - accur
Epoch 8/15
100/100 [=====] - 67s 660ms/step - loss: 1.5791 - accur
Epoch 9/15
100/100 [=====] - 67s 668ms/step - loss: 1.5556 - accur
Epoch 10/15
100/100 [=====] - 67s 669ms/step - loss: 1.5188 - accur
Epoch 11/15
100/100 [=====] - 70s 698ms/step - loss: 1.4477 - accur
Epoch 12/15
100/100 [=====] - 77s 761ms/step - loss: 1.3012 - accur
Epoch 13/15
100/100 [=====] - 76s 755ms/step - loss: 1.0356 - accur
Epoch 14/15
100/100 [=====] - 74s 734ms/step - loss: 0.7531 - accur
Epoch 15/15
100/100 [=====] - 76s 752ms/step - loss: 0.4282 - accur

```

Wow! 0.8627% accuracy! Don't get your hopes up, way too much overfitting:

tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 50 batches). You may need to use the `repeat()` function when building your dataset.

```

test_datagen = ImageDataGenerator(rescale=1./255)
test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    shuffle=False
)

test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test loss: {test_loss}, Test accuracy: {test_accuracy}")

```

Found 707 images belonging to 5 classes.

```

23/23 [=====] - 11s 485ms/step - loss: 2.7713 - accurac
Test loss: 2.7713053226470947, Test accuracy: 0.23620933294296265

```

```
train_count = sum(len(files) for _, _, files in os.walk(train_dir))
valid_count = sum(len(files) for _, _, files in os.walk(valid_dir))
test_count = sum(len(files) for _, _, files in os.walk(test_dir))

print(f"Training Images: {train_count}")
print(f"Validation Images: {valid_count}")
print(f"Test Images: {test_count}")
```

```
Training Images: 4951
Validation Images: 1415
Test Images: 707
```

```
batch_size = 32
```

```
train_steps_per_epoch = math.ceil(4951 / batch_size)
validation_steps = math.ceil(1415 / batch_size)
```

```
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=batch_size,
    class_mode='categorical'
)
```

```
validation_generator = valid_datagen.flow_from_directory(
    valid_dir,
    target_size=(150, 150),
    batch_size=batch_size,
    class_mode='categorical'
)
```

```
history = model.fit(
    train_generator,
    steps_per_epoch=train_steps_per_epoch,
    epochs=15,
    validation_data=validation_generator,
    validation_steps=validation_steps
)
```

Found 4951 images belonging to 5 classes.

Found 1415 images belonging to 5 classes.

Epoch 1/15

155/155 [=====] - 123s 776ms/step - loss: 0.2515 - accu

Epoch 2/15

155/155 [=====] - 122s 786ms/step - loss: 0.0451 - accu

Epoch 3/15

155/155 [=====] - 121s 777ms/step - loss: 0.0125 - accu

Epoch 4/15

155/155 [=====] - 128s 827ms/step - loss: 0.0015 - accu

Epoch 5/15

155/155 [=====] - 130s 837ms/step - loss: 5.5952e-04 -

Epoch 6/15

155/155 [=====] - ETA: 0s - loss: 3.7106e-04 - accuracy

**KeyboardInterrupt** Traceback (most recent call last)

<ipython-input-261-b7a1968d4ce6> in <cell line: 22>()

20

21 # Train the model

----> 22 history = model.fit(

23 train\_generator,

24 steps\_per\_epoch=train\_steps\_per\_epoch,

⏏ 13 frames

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/execute.py in  
quick\_execute(op\_name, num\_outputs, inputs, attrs, ctx, name)

51 try:

52 ctx.ensure\_initialized()

----> 53 tensors = pywrap\_tfe.TFE\_Py\_Execute(ctx.\_handle, device\_name,  
op\_name,

54 inputs, attrs, num\_outputs)

55 except core.\_NotOkStatusException as e:

**KeyboardInterrupt:**

No changes in overfitting ^

## ✓ Model 2 (Pooling and Early Stopping)

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, B
from tensorflow.keras.regularizers import l2
from tensorflow.keras.callbacks import EarlyStopping

```

```

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(2, 2),
    Dropout(0.2),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    BatchNormalization(),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Dropout(0.3),
    Flatten(),
    Dense(512, activation='relu', kernel_regularizer=l2(0.01)),
    Dropout(0.5),
    Dense(5, activation='softmax')
])

```

```

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']

```

```

early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=

```

```

history = model.fit(
    train_generator,
    steps_per_epoch=155,
    epochs=10,
    validation_data=validation_generator,
    validation_steps=45,
    callbacks=[early_stopping]
)

```

Epoch 1/10

155/155 [=====] - 117s 743ms/step - loss: 9.9326 - accu

Epoch 2/10

155/155 [=====] - 112s 724ms/step - loss: 2.6480 - accu

Epoch 3/10

155/155 [=====] - 111s 718ms/step - loss: 1.8144 - accu

Epoch 4/10

155/155 [=====] - 113s 726ms/step - loss: 1.6517 - accu

Epoch 5/10

155/155 [=====] - 111s 714ms/step - loss: 1.6185 - accu

Epoch 6/10

155/155 [=====] - 110s 710ms/step - loss: 1.6132 - accu

Epoch 7/10

155/155 [=====] - 109s 704ms/step - loss: 1.6152 - accu

Epoch 8/10

155/155 [=====] - 111s 717ms/step - loss: 1.6231 - accu

Epoch 9/10

155/155 [=====] - 116s 744ms/step - loss: 1.6194 - accu

Epoch 10/10

155/155 [=====] - 129s 833ms/step - loss: 1.6183 - accu

## ✓ Model 3: Adding Batch Normalization

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(2, 2),
    BatchNormalization(),

    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    BatchNormalization(),

    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Dropout(0.5),
    BatchNormalization(),

    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(5, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(
    train_generator,
    steps_per_epoch=155,
    epochs=10,
    validation_data=validation_generator,
    validation_steps=45,
    callbacks=[early_stopping]
)

```

Epoch 1/10

155/155 [=====] - 132s 841ms/step - loss: 2.2808 - accu

Epoch 2/10

155/155 [=====] - 110s 707ms/step - loss: 1.6092 - accu

Epoch 3/10

155/155 [=====] - 112s 719ms/step - loss: 1.5939 - accu

Epoch 4/10

155/155 [=====] - 113s 729ms/step - loss: 1.5922 - accu

Epoch 5/10

```
155/155 [=====] - 113s 731ms/step - loss: 1.5906 - accu
Epoch 6/10
155/155 [=====] - 110s 709ms/step - loss: 1.5894 - accu
Epoch 7/10
155/155 [=====] - 109s 703ms/step - loss: 1.5906 - accu
Epoch 8/10
155/155 [=====] - 108s 695ms/step - loss: 1.5798 - accu
```

```
test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}")
```

```
23/23 [=====] - 8s 363ms/step - loss: 1.5875 - accuracy
Test Loss: 1.587489366531372, Test Accuracy: 0.21074964106082916
```

These base models aren't really getting us anywhere. There were many other variations of hyperparameters that I tried but not above as I simply just changed them within the current models and just ran them again.

- ✓ Attempt 2 Let's Try Transfer Learning
- ✓ Chart of Different Models



## Available models

Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.713	0.901	138,357,544	23
VGG19	549 MB	0.713	0.900	143,667,240	26
ResNet50	98 MB	0.749	0.921	25,636,712	-
ResNet101	171 MB	0.764	0.928	44,707,176	-
ResNet152	232 MB	0.766	0.931	60,419,944	-
ResNet50V2	98 MB	0.760	0.930	25,613,800	-
ResNet101V2	171 MB	0.772	0.938	44,675,560	-
ResNet152V2	232 MB	0.780	0.942	60,380,648	-
InceptionV3	92 MB	0.779	0.937	23,851,784	159
InceptionResNetV2	215 MB	0.803	0.953	55,873,736	572
MobileNet	16 MB	0.704	0.895	4,253,864	88
MobileNetV2	14 MB	0.713	0.901	3,538,984	88
DenseNet121	33 MB	0.750	0.923	8,062,504	121
DenseNet169	57 MB	0.762	0.932	14,307,880	169
DenseNet201	80 MB	0.773	0.936	20,242,984	201
NASNetMobile	23 MB	0.744	0.919	5,326,716	-
NASNetLarge	343 MB	0.825	0.960	88,949,818	-
EfficientNetB0	29 MB	-	-	5,330,571	-
EfficientNetB1	31 MB	-	-	7,856,239	-
EfficientNetB2	36 MB	-	-	9,177,569	-
EfficientNetB3	48 MB	-	-	12,320,535	-
EfficientNetB4	75 MB	-	-	19,466,823	-
EfficientNetB5	118 MB	-	-	30,562,527	-
EfficientNetB6	166 MB	-	-	43,265,143	-
EfficientNetB7	256 MB	-	-	66,658,687	-

This chart shows us many options that we can proceed with, for the sake of time and available hardware we decided to test out 2, VGG16 and ResNet50.

Chart from: <https://medium.com/@blant.jesse/transfer-learning-neur-12df2f55b601>

## ▼ ResNet50

```

from tensorflow.keras.applications import ResNet50

train_datagen = ImageDataGenerator(
    preprocessing_function=tf.keras.applications.resnet.preprocess_input,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)

valid_datagen = ImageDataGenerator(
    preprocessing_function=tf.keras.applications.resnet.preprocess_input
)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)

validation_generator = valid_datagen.flow_from_directory(
    valid_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)

base_model = ResNet50(include_top=False, weights='imagenet', input_shape=(224, 224,

for layer in base_model.layers:
    layer.trainable = False

model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(1024, activation='relu'),
    Dense(5, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

Found 4951 images belonging to 5 classes.
Found 1415 images belonging to 5 classes.
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/94765736/94765736 [=====] - 0s 0us/step

model.summary()

```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0
dense_18 (Dense)	(None, 1024)	2098176
dense_19 (Dense)	(None, 5)	5125
Total params: 25691013 (98.00 MB)		
Trainable params: 2103301 (8.02 MB)		
Non-trainable params: 23587712 (89.98 MB)		

```
history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // train_generator.batch_size,
    epochs=10,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // validation_generator.batch_size
)
```

```
Epoch 1/10
154/154 [=====] - 354s 2s/step - loss: 1.8487 - accurac
Epoch 2/10
154/154 [=====] - 344s 2s/step - loss: 1.5682 - accurac
Epoch 3/10
154/154 [=====] - 354s 2s/step - loss: 1.5639 - accurac
Epoch 4/10
154/154 [=====] - 351s 2s/step - loss: 1.5534 - accurac
Epoch 5/10
154/154 [=====] - 351s 2s/step - loss: 1.5445 - accurac
Epoch 6/10
154/154 [=====] - 342s 2s/step - loss: 1.5367 - accurac
Epoch 7/10
154/154 [=====] - 340s 2s/step - loss: 1.5338 - accurac
Epoch 8/10
154/154 [=====] - 337s 2s/step - loss: 1.5298 - accurac
Epoch 9/10
154/154 [=====] - 339s 2s/step - loss: 1.5233 - accurac
Epoch 10/10
154/154 [=====] - 339s 2s/step - loss: 1.5109 - accurac
```

```
test_datagen = ImageDataGenerator(preprocessing_function=tf.keras.applications.resne
test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)
```

```
test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}")
```

Found 707 images belonging to 5 classes.

23/23 [=====] - 38s 2s/step - loss: 1.5787 - accuracy:

Test Loss: 1.5786525011062622, Test Accuracy: 0.3055162727832794

This works better than before, accuracy went up with not much change in loss in comparison to the previous models we tried. Now onto VGG16.

## ✓ VGG16

```
from tensorflow.keras.applications import VGG16

train_datagen = ImageDataGenerator(
    preprocessing_function=tf.keras.applications.vgg16.preprocess_input,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

valid_datagen = ImageDataGenerator(
    preprocessing_function=tf.keras.applications.vgg16.preprocess_input
)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)

validation_generator = valid_datagen.flow_from_directory(
    valid_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical'
)

base_model = VGG16(include_top=False, weights='imagenet', input_shape=(224, 224, 3))

for layer in base_model.layers:
    layer.trainable = False

model = Sequential([
    base_model,
    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(5, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

Found 4951 images belonging to 5 classes.
Found 1415 images belonging to 5 classes.
```

```
model.summary()
```

```
Model: "sequential_15"
```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_14 (Flatten)	(None, 25088)	0
dense_30 (Dense)	(None, 512)	12845568
dropout_15 (Dropout)	(None, 512)	0
dense_31 (Dense)	(None, 5)	2565

---

Total params: 27562821 (105.14 MB)  
Trainable params: 12848133 (49.01 MB)  
Non-trainable params: 14714688 (56.13 MB)

---

```
history = model.fit(  
    train_generator,  
    steps_per_epoch=train_generator.samples // train_generator.batch_size,  
    epochs=10,  
    validation_data=validation_generator,  
    validation_steps=validation_generator.samples // validation_generator.batch_size  
)
```

```

Epoch 1/10
154/154 [=====] - 964s 6s/step - loss: 10.1694 - accura
Epoch 2/10
154/154 [=====] - 927s 6s/step - loss: 1.6332 - accurac
Epoch 3/10
154/154 [=====] - 997s 6s/step - loss: 1.6144 - accurac
Epoch 4/10
81/154 [=====>.....] - ETA: 5:47 - loss: 1.6150 - accuracy:

```

**KeyboardInterrupt** Traceback (most recent call last)

<ipython-input-283-9499a636fc1d> in <cell line: 1>()

```

----> 1 history = model.fit(
      2     train_generator,
      3     steps_per_epoch=train_generator.samples //
train_generator.batch_size,
      4     epochs=10,
      5     validation_data=validation_generator,

```

⏮ 10 frames ⏭

```

/usr/local/lib/python3.10/dist-packages/tensorflow/python/eager/execute.py in
quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)

```

```

    51     try:
    52         ctx.ensure_initialized()
----> 53         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name,
op_name,
    54                                     inputs, attrs, num_outputs)
    55     except core._NotOkStatusException as e:

```

**KeyboardInterrupt:**

We stopped running this as we have already run the whole thing in a previous notebook and it is taking too long but here are the results from VGG16 from the previous notebook. Given to the time constraint we chose to continue along with ResNet50



```

Epoch 4/20
124/124 [=====] - 238s 2s/step - loss: 1.4108 - accuracy: 0.2831 - val_loss: 1.9750 - val_accuracy: 0.2311
Epoch 5/20
124/124 [=====] - 251s 2s/step - loss: 1.4062 - accuracy: 0.2930 - val_loss: 2.0215 - val_accuracy: 0.2332
Epoch 6/20
124/124 [=====] - 248s 2s/step - loss: 1.4074 - accuracy: 0.2864 - val_loss: 1.9368 - val_accuracy: 0.2332
Epoch 7/20
124/124 [=====] - 248s 2s/step - loss: 1.4069 - accuracy: 0.2836 - val_loss: 1.8395 - val_accuracy: 0.2254
Epoch 8/20
124/124 [=====] - 250s 2s/step - loss: 1.4062 - accuracy: 0.2975 - val_loss: 2.1161 - val_accuracy: 0.2332
Epoch 9/20
124/124 [=====] - 251s 2s/step - loss: 1.4061 - accuracy: 0.2811 - val_loss: 2.0003 - val_accuracy: 0.2311
Epoch 10/20
124/124 [=====] - 251s 2s/step - loss: 1.4116 - accuracy: 0.2755 - val_loss: 1.8909 - val_accuracy: 0.2311
Epoch 11/20
124/124 [=====] - 247s 2s/step - loss: 1.4069 - accuracy: 0.2879 - val_loss: 1.9056 - val_accuracy: 0.2332
Epoch 12/20
124/124 [=====] - 249s 2s/step - loss: 1.4093 - accuracy: 0.2728 - val_loss: 2.0957 - val_accuracy: 0.2332
Epoch 13/20
124/124 [=====] - 241s 2s/step - loss: 1.4066 - accuracy: 0.2839 - val_loss: 2.0118 - val_accuracy: 0.2311
Epoch 14/20
124/124 [=====] - 249s 2s/step - loss: 1.4057 - accuracy: 0.2877 - val_loss: 2.0179 - val_accuracy: 0.2311
Epoch 15/20
124/124 [=====] - 248s 2s/step - loss: 1.4063 - accuracy: 0.2861 - val_loss: 2.0415 - val_accuracy: 0.2311
Epoch 16/20
124/124 [=====] - 250s 2s/step - loss: 1.4068 - accuracy: 0.2882 - val_loss: 1.9349 - val_accuracy: 0.2311
Epoch 17/20
124/124 [=====] - 252s 2s/step - loss: 1.4060 - accuracy: 0.2745 - val_loss: 1.9207 - val_accuracy: 0.2311
Epoch 18/20
124/124 [=====] - 250s 2s/step - loss: 1.4046 - accuracy: 0.2831 - val_loss: 1.9695 - val_accuracy: 0.2332
Epoch 19/20
124/124 [=====] - 236s 2s/step - loss: 1.4057 - accuracy: 0.2776 - val_loss: 1.9698 - val_accuracy: 0.2254
Epoch 20/20

```

As you can see there is no benefit to using VGG16 over ResNet50 so we proceeded with ResNet50.

## ✓ Attempt 3 Redistributing the Data

```
!ls drive/MyDrive/ds340/project/o
```

```
test train valid
```

```

import os
import shutil
from sklearn.model_selection import train_test_split

def create_dataset_copy(original_dir, copy_dir):
    """Copies the dataset from the original directory to a new directory."""
    if not os.path.exists(copy_dir):
        shutil.copytree(original_dir, copy_dir)
    else:
        print(f"Copy directory {copy_dir} already exists.")

original_base_dir = '/content/drive/MyDrive/ds340/project/data1'
copy_base_dir = '/content/drive/MyDrive/ds340/project/data_copy'

create_dataset_copy(original_base_dir, copy_base_dir)

def consolidate_images(base_dir, classes, set_names):
    """This function now works on the dataset copy."""
    consolidated_dir = os.path.join(base_dir, 'consolidated')
    os.makedirs(consolidated_dir, exist_ok=True)

    for cls in classes:
        class_dir = os.path.join(consolidated_dir, cls)
        os.makedirs(class_dir, exist_ok=True)
        for set_name in set_names:
            src_dir = os.path.join(base_dir, set_name, cls)
            if os.path.exists(src_dir):
                for file in os.listdir(src_dir):
                    src_file_path = os.path.join(src_dir, file)
                    dst_file_path = os.path.join(class_dir, file)
                    if not os.path.exists(dst_file_path):
                        shutil.copy(src_file_path, dst_file_path) # Change from mov

classes = ['Folding Knife', 'Multi-tool Knife', 'Straight Knife', 'Utility Knife', '
set_names = ['train', 'valid', 'test']
consolidate_images(copy_base_dir, classes, set_names)

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-290-3110820b0a69> in <cell line: 55>()
    53         move_files(test_files, os.path.join(base_dir, 'test', cls))
    54
--> 55 redistribute_images(copy_base_dir, train_ratio=0.70, val_ratio=0.15)
    56
    57 # Example to count images in the copied directory

<ipython-input-290-3110820b0a69> in redistribute_images(base_dir, train_ratio,
val_ratio)
    49         train_files, val_files = train_test_split(train_files,
test_size=val_ratio/(train_ratio+val_ratio), random_state=42)
    50
--> 51         move_files(train_files, os.path.join(base_dir, 'train', cls))
    52         move_files(val_files, os.path.join(base_dir, 'valid', cls))
    53         move_files(test_files, os.path.join(base_dir, 'test', cls))

NameError: name 'move_files' is not defined

```

```
import os
```

```

def count_images(directory):
    """ Count the number of images in each class directory. """
    class_counts = {}
    for class_dir in os.listdir(directory):
        class_path = os.path.join(directory, class_dir)
        if os.path.isdir(class_path): # Ensure it's a directory
            count = len(os.listdir(class_path))
            class_counts[class_dir] = count
    return class_counts

```

```

train_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_
val_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_co
test_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_c

```

```

print("Training set counts:", train_counts)
print("Validation set counts:", val_counts)
print("Test set counts:", test_counts)

```

```

Training set counts: {'Utility Knife': 1110, 'Folding Knife': 1103, 'Multi-tool
Validation set counts: {'Utility Knife': 327, 'Folding Knife': 319, 'Multi-tool
Test set counts: {'Utility Knife': 177, 'Folding Knife': 155, 'Multi-tool Knife'

```

```

import os
import shutil

def consolidate_images(base_dir, classes, set_names):
    consolidated_dir = os.path.join(base_dir, 'consolidated')
    os.makedirs(consolidated_dir, exist_ok=True)

    for cls in classes:
        class_dir = os.path.join(consolidated_dir, cls)
        os.makedirs(class_dir, exist_ok=True)

        for set_name in set_names:
            src_dir = os.path.join(base_dir, set_name, cls)
            if os.path.exists(src_dir):
                for file in os.listdir(src_dir):
                    src_file_path = os.path.join(src_dir, file)
                    dst_file_path = os.path.join(class_dir, file)
                    if not os.path.exists(dst_file_path):
                        shutil.move(src_file_path, dst_file_path)

base_dir = '/content/drive/MyDrive/ds340/project/data_copy'
classes = ['Folding Knife', 'Multi-tool Knife', 'Straight Knife', 'Utility Knife', ' '
set_names = ['train', 'valid', 'test']
consolidate_images(base_dir, classes, set_names)

from sklearn.model_selection import train_test_split

def redistribute_images(base_dir, train_ratio=0.60, val_ratio=0.20):
    consolidated_dir = os.path.join(base_dir, 'consolidated')
    classes = os.listdir(consolidated_dir)

    for cls in classes:
        class_dir = os.path.join(consolidated_dir, cls)
        files = [os.path.join(class_dir, f) for f in os.listdir(class_dir)]

        train_files, test_files = train_test_split(files, test_size=1-train_ratio-va
        train_files, val_files = train_test_split(train_files, test_size=val_ratio/(

    def move_files(files, target_dir):
        os.makedirs(target_dir, exist_ok=True)
        for f in files:
            shutil.move(f, target_dir)

    move_files(train_files, os.path.join(base_dir, 'train', cls))
    move_files(val_files, os.path.join(base_dir, 'valid', cls))
    move_files(test_files, os.path.join(base_dir, 'test', cls))

redistribute_images(base_dir, train_ratio=0.70, val_ratio=0.15)

```

```

import os

def count_images(directory):
    """ Count the number of images in each class directory. """
    class_counts = {}
    for class_dir in os.listdir(directory):
        class_path = os.path.join(directory, class_dir)
        if os.path.isdir(class_path): # Ensure it's a directory
            count = len(os.listdir(class_path))
            class_counts[class_dir] = count
    return class_counts

train_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_val_
val_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_co
test_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_c

print("Training set counts:", train_counts)
print("Validation set counts:", val_counts)
print("Test set counts:", test_counts)

```

```

Training set counts: {'Utility Knife': 1129, 'Folding Knife': 1102, 'Multi-tool
Validation set counts: {'Utility Knife': 242, 'Folding Knife': 238, 'Multi-tool
Test set counts: {'Utility Knife': 243, 'Folding Knife': 237, 'Multi-tool Knife'

```

### Oversample Miniority Class - Data Augment

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    fill_mode='nearest'
)

```

```

import os
import tensorflow as tf

base_dir = '/content/drive/MyDrive/ds340/project/data_copy'
class_name = 'Straight Knife'
train_dir = os.path.join(base_dir, 'train', class_name)
valid_dir = os.path.join(base_dir, 'valid', class_name)
test_dir = os.path.join(base_dir, 'test', class_name)

def augment_images(directory, num_augmented_per_image=5):
    for filename in os.listdir(directory):
        if filename.endswith('.jpg') or filename.endswith('.png'):
            image_path = os.path.join(directory, filename)
            image = tf.keras.preprocessing.image.load_img(image_path)
            image = image.resize((150, 150))
            x = tf.keras.preprocessing.image.img_to_array(image)
            x = x.reshape((1,) + x.shape)

            i = 0
            for batch in datagen.flow(x, batch_size=1, save_to_dir=directory, save_p
                i += 1
                if i >= num_augmented_per_image:
                    break

augment_images(train_dir)
augment_images(valid_dir)
augment_images(test_dir)

import os

def count_images(directory):
    """ Count the number of images in each class directory. """
    class_counts = {}
    for class_dir in os.listdir(directory):
        class_path = os.path.join(directory, class_dir)
        if os.path.isdir(class_path): # Ensure it's a directory
            count = len(os.listdir(class_path))
            class_counts[class_dir] = count
    return class_counts

train_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_
val_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_co
test_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_c

print("Training set counts:", train_counts)
print("Validation set counts:", val_counts)
print("Test set counts:", test_counts)

```

Training set counts: {'Utility Knife': 1129, 'Folding Knife': 1102, 'Multi-tool'  
 Validation set counts: {'Utility Knife': 242, 'Folding Knife': 238, 'Multi-tool'  
 Test set counts: {'Utility Knife': 243, 'Folding Knife': 237, 'Multi-tool Knife'

```
import os
import random

def undersample_directory(directory, target_count):
    """ Randomly remove files from a directory to reduce to target_count. """
    files = [os.path.join(directory, f) for f in os.listdir(directory) if f.endswith(
    current_count = len(files)
    if current_count <= target_count:
        print(f"No need to remove files from {directory}, count is already at or bel
        return

    remove_count = current_count - target_count
    files_to_remove = random.sample(files, remove_count)

    for file in files_to_remove:
        os.remove(file)
    print(f"Removed {remove_count} files from {directory}")

base_dir = '/content/drive/MyDrive/ds340/project/data_copy'
train_uk_dir = os.path.join(base_dir, 'train', 'Utility Knife')
valid_uk_dir = os.path.join(base_dir, 'valid', 'Utility Knife')
test_uk_dir = os.path.join(base_dir, 'test', 'Utility Knife')

target_train_count = 700
target_valid_count = 150
target_test_count = 150

undersample_directory(train_uk_dir, target_train_count)
undersample_directory(valid_uk_dir, target_valid_count)
undersample_directory(test_uk_dir, target_test_count)
```

Removed 429 files from /content/drive/MyDrive/ds340/project/data\_copy/train/Util  
 Removed 92 files from /content/drive/MyDrive/ds340/project/data\_copy/valid/Utili  
 Removed 93 files from /content/drive/MyDrive/ds340/project/data\_copy/test/Utilit

```
import os

def count_images(directory):
    """ Count the number of images in each class directory. """
    class_counts = {}
    for class_dir in os.listdir(directory):
        class_path = os.path.join(directory, class_dir)
        if os.path.isdir(class_path): # Ensure it's a directory
            count = len(os.listdir(class_path))
            class_counts[class_dir] = count
    return class_counts

train_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_val_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_co
test_counts = count_images(os.path.join('/content/drive/MyDrive/ds340/project/data_c

print("Training set counts:", train_counts)
print("Validation set counts:", val_counts)
print("Test set counts:", test_counts)
```

```
Training set counts: {'Utility Knife': 700, 'Folding Knife': 1102, 'Multi-tool K
Validation set counts: {'Utility Knife': 150, 'Folding Knife': 238, 'Multi-tool
Test set counts: {'Utility Knife': 150, 'Folding Knife': 237, 'Multi-tool Knife'
```

```
import tensorflow as tf
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator

def build_model(learning_rate, dropout_rate, optimizer):
    base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 2

    for layer in base_model.layers:
        layer.trainable = False

    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    x = Dropout(dropout_rate)(x)
    predictions = Dense(5, activation='softmax')(x)

    model = Model(inputs=base_model.input, outputs=predictions)

    model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['ac

    return model
```



```
learning_rates = [0.001, 0.0001, 0.00001]  
dropout_rates = [0.3, 0.5, 0.7]  
optimizers = [tf.keras.optimizers.Adam, tf.keras.optimizers.RMSprop, tf.keras.optimi
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam

train_dir = '/content/drive/MyDrive/ds340/project/data_copy/train'
val_dir = '/content/drive/MyDrive/ds340/project/data_copy/valid'
test_dir = '/content/drive/MyDrive/ds340/project/data_copy/test'

# Define image size and batch size
img_height, img_width = 150, 150 # Adjust these dimensions to your specific dataset
batch_size = 32

# Prepare the data generators
train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
val_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
test_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical')

validation_generator = val_datagen.flow_from_directory(
    val_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical')

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical')

train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)
```

```
Found 6990 images belonging to 5 classes.
Found 1561 images belonging to 5 classes.
Found 1561 images belonging to 5 classes.
```

```
for learning_rate in learning_rates:
    for dropout_rate in dropout_rates:
        for optimizer_class in optimizers:
            optimizer = optimizer_class(learning_rate=learning_rate)
            model = build_model(learning_rate, dropout_rate, optimizer)

            history = model.fit(train_generator,
                                steps_per_epoch=train_generator.samples // train_gen
                                epochs=10,
                                validation_data=validation_generator,
                                validation_steps=validation_generator.samples // val

            val_loss, val_acc = model.evaluate(validation_generator,
                                                steps=validation_generator.samples //

            print(f'Learning Rate: {learning_rate}, Dropout Rate: {dropout_rate}, Op
            print(f'Validation Loss: {val_loss}, Validation Accuracy: {val_acc}')
            print('-' * 50)
```

Epoch 1/10

196/196 [=====] - 47s 224ms/step - loss: 2.3437 - accur

Epoch 2/10

196/196 [=====] - 41s 211ms/step - loss: 2.1287 - accur

Epoch 3/10

196/196 [=====] - 43s 221ms/step - loss: 1.9778 - accur

Epoch 4/10

196/196 [=====] - 43s 222ms/step - loss: 1.8564 - accur

Epoch 5/10

196/196 [=====] - 42s 215ms/step - loss: 1.7765 - accur

Epoch 6/10

196/196 [=====] - 42s 216ms/step - loss: 1.7261 - accur

Epoch 7/10

196/196 [=====] - 42s 216ms/step - loss: 1.6948 - accur

Epoch 8/10

196/196 [=====] - 42s 214ms/step - loss: 1.6814 - accur

Epoch 9/10

196/196 [=====] - 42s 215ms/step - loss: 1.6621 - accur

Epoch 10/10

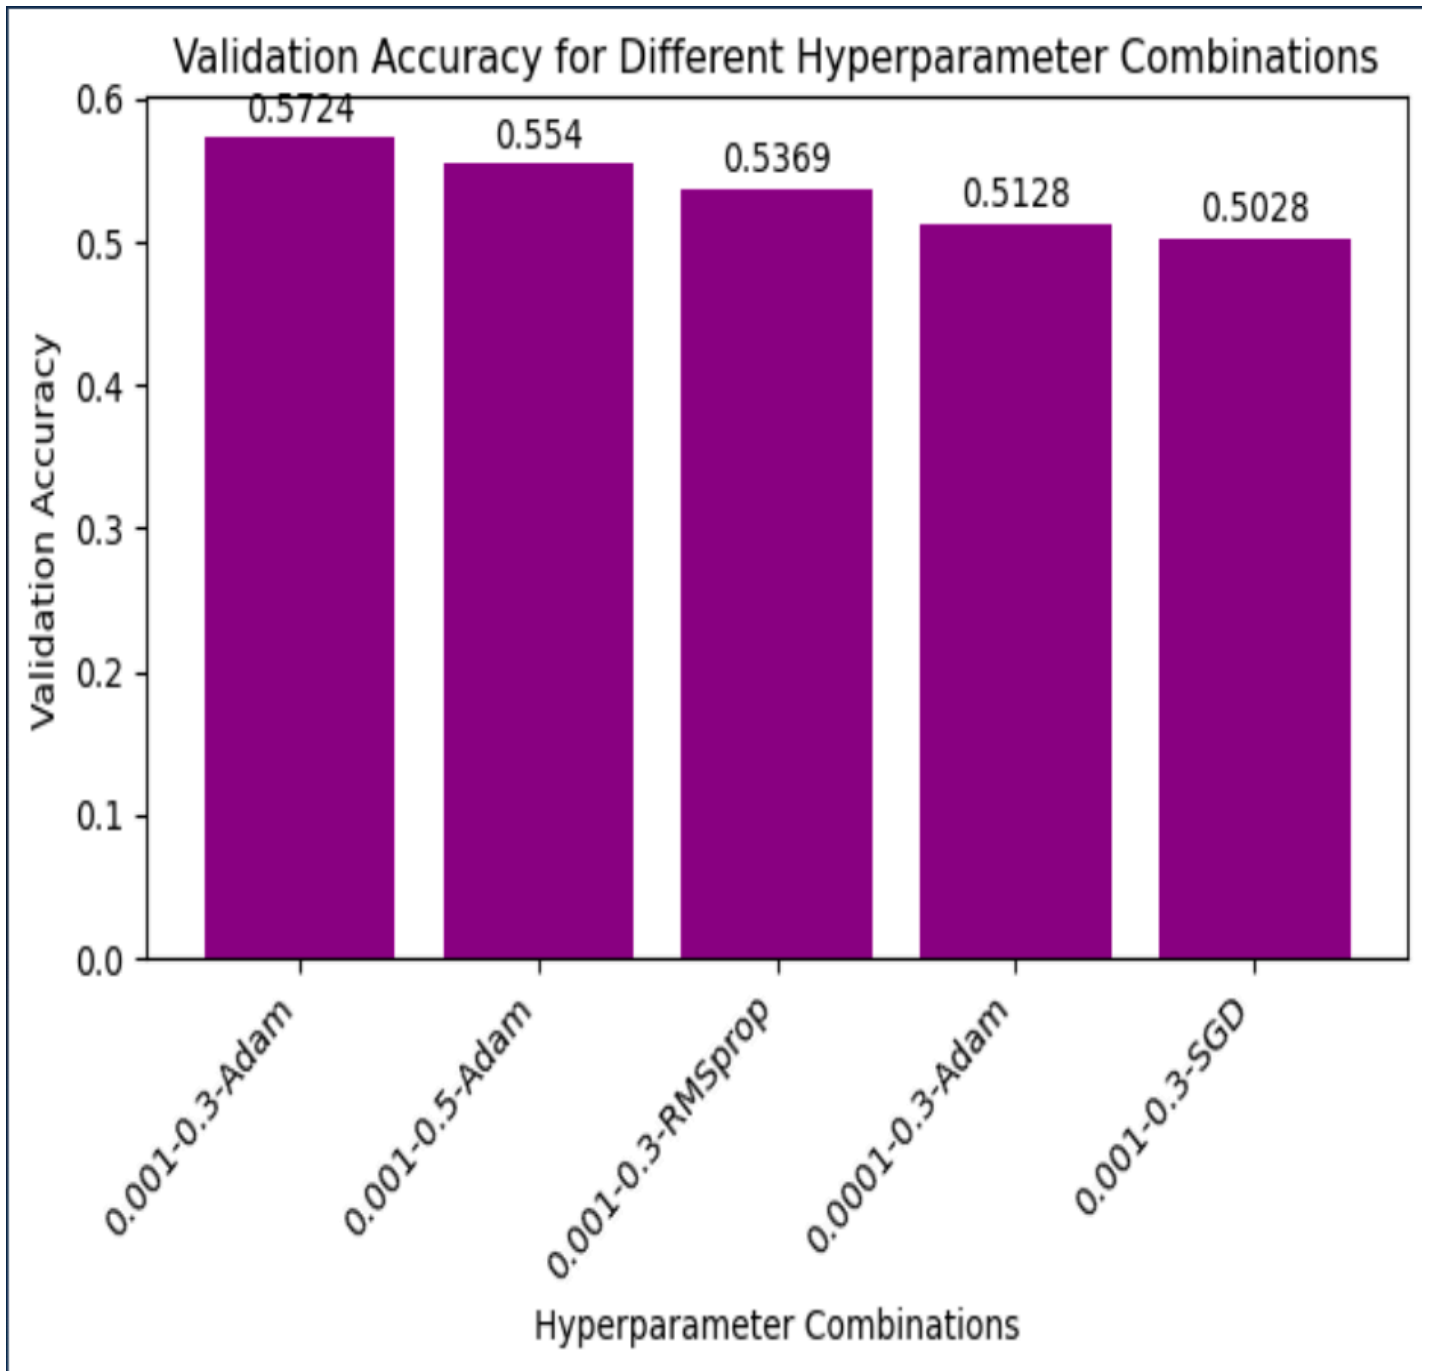
196/196 [=====] - 43s 220ms/step - loss: 1.6378 - accur


44/44 [=====] - 8s 176ms/step - loss: 1.3672 - accuracy

Learning Rate: 1e-05, Dropout Rate: 0.7, Optimizer: SGD

Validation Loss: 1.367194414138794, Validation Accuracy: 0.5028409361839294

-----





- KEY:
- Config 1: Learning Rate: 0.001, Dropout Rate: 0.3, Optimizer: Adam
  - Config 2: Learning Rate: 0.001, Dropout Rate: 0.3, Optimizer: RMSprop
  - Config 3: Learning Rate: 0.001, Dropout Rate: 0.3, Optimizer: SGD
  - Config 4: Learning Rate: 0.001, Dropout Rate: 0.5, Optimizer: Adam
  - Config 5: Learning Rate: 0.0001, Dropout Rate: 0.3, Optimizer: Adam
  - Config 6: Learning Rate: 0.001, Dropout Rate: 0.5, Optimizer: RMSprop
  - Config 7: Learning Rate: 0.001, Dropout Rate: 0.7, Optimizer: Adam
  - Config 8: Learning Rate: 0.0001, Dropout Rate: 0.5, Optimizer: Adam
  - Config 9: Learning Rate: 0.0001, Dropout Rate: 0.3, Optimizer: RMSprop
  - Config 10: Learning Rate: 0.0001, Dropout Rate: 0.7, Optimizer: Adam

